

CS-417

COMPUTER SYSTEMS MODELING

Spring Semester 2020

Batch: 2016-17
(LECTURE # 4)

FAKHRA AFTAB
LECTURER

DEPARTMENT OF COMPUTER & INFORMATION SYSTEMS ENGINEERING
NED UNIVERSITY OF ENGINEERING & TECHNOLOGY



Recap of Lecture # 3

Amdahl's Law

Means vs. End based Metric

Performance Measures

Review of Computer Systems Performance Evaluation



Chapter # 1 (Cont'd)

COMPUTER SYSTEMS PERFORMANCE MODELING AND EVALUATION



Benchmarking

Benchmarking (in computing) is the act of running a set of computer programs, or other operations, in order to assess the relative performance of an object in the computer system.

Benchmark: elaborately-designed benchmarking programs.

- Usually associated with assessing performance of hardware.
- But the technique is also applicable to software.



Purpose of Benchmarking

- Provides a method of comparing the performance of various systems (or subsystems) across different architectures.
- Benchmarks are designed to mimic a particular type of workload on a component or system.
- It also helps to choose effective system upgrades and to identify potential bottlenecks.



Types of Benchmarks

i) Addition Instruction

- Historically, the performance of computer system was synonymous with that of the processor.
- Addition was the most frequent instruction.
- The machine with the fastest addition operation would produce the best overall performance when executing any application program.



Types of Benchmarks (Cont'd)

ii) Instruction Mixes

- Instruction sets in which each instruction take only the minimum number of cycles required to complete its particular operation were designed.
- Execution time of a single instruction was no longer adequate to summarize performance.
- Gibson proposed instruction mix as benchmarking.
- The basic idea was to categorize all of the instructions into different classes.
- Number of instructions of each class executed by the particular collection of programs was then used to form a weighted average.



Example: Gibson Instruction Mix

1.	Load and Store	13.2
2.	Fixed-Point Add/Sub	6.1
3.	Compares	3.8
4.	Branches	16.6
5.	Float Add/Sub	6.9
6.	Float Multiply	3.8
7.	Float Divide	1.5
8.	Fixed-Point Multiply	0.6
9.	Fixed-Point Divide	0.2
10.	Shifting	4.4
11.	Logical And/Or	1.6
12.	Instructions not using regs	5.3
13.	Indexing	<u>18.0</u>
Total		100

1959,
IBM 650
IBM 704



Types of Benchmarks (Cont'd)

iii) Kernel

- Small benchmark codes designed to measure basic architectural features of parallel machines
 - Normally abstracted from actual program
 - Easy to port to many different computer systems.
 - Most kernels are processing kernels – the only focus is on processor's performance

Popular kernels

- Livermore loops consists of 24 do loops, some of which can be vectorized, and some of which cannot
- Linpack benchmark (contains basic linear algebra subroutine written in FORTRAN language)
- Tower of Hanoi, 8-Queen Puzzle Problem, Sorting etc.



Types of Benchmarks (Cont'd)

iv) Component Benchmark / Micro-benchmark

- Consists of a relatively small specific piece of code to measure performance of computer's specific portion/component.
- Example: A small program written to test only the processor-memory interface, the input/output subsystem, or the FP execution unit independent of the other components of the system.
- Determine whether the performance capabilities of all the components within a system are balanced.
- Provide values for important parameters for a simulation of the system.
- Requires programmer to have a deep understanding of system being tested.



Types of Benchmarks (Cont'd)

Some more examples:

- The most important component in a particular case depends on the individual's usage patterns.
 - A gamer seeking the best possible frame rates will probably be better served by a faster GPU than by more memory.
 - A casual user seeking more responsive system may benefit by upgrading a slow hard drive to fast solid-state drive.
 - It is necessary to tailor the suite of benchmarks according to user's specific needs and then weigh the individual test results accordingly.



Types of Benchmarks (Cont'd)

v) Synthetic Benchmarks

- Presents a mix of computations, system calls and I/O requests.
- Artificial programs that mimic a real program execution environment by using statistical data about real high-level language programs.
- Can be developed quickly and do not rely on real data.
- Some examples of the relevant statistical data are:
 - Fraction of statements of each type (assignment, conditionals, for-loops, etc.)
 - Fraction of variables of each type (integer, real, character, etc.) and locality (local or global)
 - Fraction of expressions with certain number and type of operators and operands



Types of Benchmarks (Cont'd)

- Procedure for programming synthetic benchmark:
 - take statistics of all types of operations from many application programs
 - get proportion of each operation
 - write program based on the proportion above
- They are too small to make statistically relevant memory or disk references.

vi) Application Benchmarks (or Real programs)

- A small subset of real-world programs that are representative of the application domain of interest.
- e.g. word processing software, tool software of CAD, user's application software (i.e. MIS)



Types of Benchmarks (Cont'd)

vii) Database benchmarks

- To measure the throughput and response times of database management systems (DBMS)

viii) Parallel benchmarks

- Used on machines with multiple cores, processors or systems consisting of multiple machines.



Application vs Synthetic benchmarks

Main advantage of synthetic benchmark over application one:

- overall application domain characteristics can be closely matched by a single program.
- Application benchmarks *give much better measure of real-world performance on* a given system,
- Synthetic benchmarks are *useful for testing individual components*, like a hard disk or networking device.



Application vs Synthetic benchmarks (Cont'd)

Since a synthetic benchmark not designed to do anything meaningful, may show following *strange characteristics*:

- 1) Since the expressions controlling conditional statements are chosen randomly,
 - a significant amount of unreachable (or *dead*) code may be created, much of which can be eliminated by a *smart* compiler.
 - Code-size becomes highly dependent on quality of compiler.
- 2) Random selection of statements may result in *unusual locality properties*, defeating well-designed paging algorithms.
- 3) The benchmark may be *small enough* to fit entirely in the cache and result in an unusually good performance.



Some Popular Benchmarks

i) Whetstone synthetic benchmark

- Based upon the characteristics of FORTRAN programs doing extensive floating-point computation.
- A single *Whetstone instruction* is defined as a mix of basic high-level operations, and the results are reported as *mega-Whetstones per second*.
- This benchmark, although still very popular represents outdated and arbitrary instruction mixes and thus
 - not very useful for machines with a high degree of internal parallelism (pipelining, vectored computation etc.),
 - particularly in conjunction with optimizing and parallelizing compilers.



Some Popular Benchmarks (Cont'd)

ii) Dhrystone synthetic benchmark

- written in C,
- designed to represent applications involving primarily integer arithmetic and string manipulation in a block-structured language.
- This benchmark is only 100 statements long, and thus would fit in the instruction cache of most systems.
- Also, since calls only two simple procedures, *strcpy()* and *strcmp()*, a compiler that replaces them by inline code can cause a very significant speed-up.
- The results are quoted as *Dhrystones/sec*, i.e. the rate at which the individual benchmark statements are executed.



Some Popular Benchmarks (Cont'd)

iii) Linpack application benchmark

- Solves a dense 100 x 100 linear system of equations using the Linpack library package.
- Two problems with this benchmark:
 - Over 80% of the time is spent in doing $A(I) = B(I) + C \times D(I)$ type of calculation, making results highly dependent on how well such a computation is handled,
 - Problem is too small.



Precautions in Benchmarking

A number of *hardware and software factors* need to be considered before running benchmarks.

- 1) Many benchmarks place *significant stress* on specific components,
 - ensure all such components are in *good working order*,
 - *properly cooled* (if necessary), and *receiving adequate power*.
- 2) When comparing machines using benchmarks, some care is needed to eliminate the effect of *extraneous factors*.
 - e.g., a highly optimizing compiler can often speedup the code *twicely* compared with an ordinary compiler.
 - Therefore, it's necessary that benchmarks for all machines be compiled using *compilers with similar features and settings*.
- 3) The details of machine configuration should be spelled out clearly, e.g., one could not claim much shorter times by just increasing the cache-size.



Precautions in Benchmarking (Cont'd)

4) Parameters for OS, applications and drivers must be satisfied to ensure accurate, repeatable benchmark results.

- Windows (and other) OSs proactively prefetch data and store numerous temporary files that could interfere with a benchmark, so it's best to clear out any temporary files and prefetch data before running a test.
- In Windows 7, you can find prefetch data in *C:\Windows\Prefetch*, and temporary files in *C:\Windows\Temp* and *C:\Users\[username]\AppData\Local\Temp*.

5) Using the correct (or latest) drivers for a component to ensure it is operating and performing optimally.

- especially true of graphics boards and motherboards /chipsets, where wrong driver can significantly worsen the system's frame rates or transfer speeds and latency.
- Finally, confirm that the OS is fully updated and patched to ensure optimal compatibility and to reflect the current, real-world OS configuration.



List of Challenges

Interpretation of benchmarking data is extraordinarily difficult.

- 1) Manufacturers commonly report only those benchmarks (or aspects of benchmarks) that show their products in the best light.
- 2) Many benchmarks focus entirely on the speed of computational performance, neglecting other important features such as *Qualities of Service*.
 - Transaction Processing Performance Council (TPPC) Benchmark specifications partially address these concerns by specifying *ACID property tests*, *database scalability rules* and *service level requirements*.



List of Challenges (Cont'd)

3) Users can have very different perceptions of performance than benchmarks may suggest.

- In particular, users appreciate *predictability* — servers that always meet or exceed service level agreements.
- Benchmarks tend to emphasize mean scores (IT perspective), rather than maximum worst-case response times (real-time computing perspective), or low standard deviations (user perspective).

4) Many server architectures degrade dramatically at high (near 100%) levels of usage.

- Vendors, in particular, tend to publish server benchmarks at continuous at about 80% usage — an unrealistic situation.



Benchmark Strategies

Most of the benchmark programs base the measure of performance on *time required* to execute benchmark. Several other strategies can be employed, however.

Specifically, the three different strategies for using a benchmark program are the following.

1) Fixed Computation Benchmarks:

Measure the time required to perform a *fixed amount of computation*.

- a fixed workload
- Example: The SPEC (Standard Performance Evaluation Corporation) CPU benchmarks
- www.spec.org



Benchmark Strategies (Cont'd)

2) Fixed Time Benchmarks:

Measure the amount of computation performed within a *fixed period of time*.

- *Argument:* users with large problems are willing to wait a fixed amount of time to obtain a solution.
- At the end of allotted execution time, total amount of computation completed, used as the measure of relative speeds of different systems.
- Example: The TPC (Transaction Processing perf. Council) benchmarks.
- www.tpc.org



Benchmark Strategies (Cont'd)

3) Variable-Computation & Variable-Time Benchmarks:

Allow both the amount of computation ~~performed~~ and the time to vary.

- ~~These types of benchmarks instead try to~~ measure some ~~other~~ aspect of performance that is a function of the computation and the execution time.
- Example: The HINT (Hierarchical INTegration) benchmark.
- HINT uses a performance measure called QUIPS (QUality Improvements Per Second).

